

Joshua Kobuskie

In [1]:

```
#INSTALL JAX AND OTHER LIBRARIES
```

```
!pip install jax
import jax
import pandas as pd
import numpy as np
import gc
import sys
```

Requirement already satisfied: jax in /opt/conda/lib/python3.7/site-packages (0.3.1)

Requirement already satisfied: numpy>=1.19 in /opt/conda/lib/python3.7/site-packages (from jax) (1.20.3)

Requirement already satisfied: opt-einsum in /opt/conda/lib/python3.7/site-packages (from jax) (3.3.0)

Requirement already satisfied: scipy>=1.2.1 in /opt/conda/lib/python3.7/site-packages (from jax) (1.7.3)

Requirement already satisfied: typing-extensions in /opt/conda/lib/python3.7/site-packages (from jax) (4.1.1)

Requirement already satisfied: absl-py in /opt/conda/lib/python3.7/site-packages (from jax) (0.15.0)

Requirement already satisfied: six in /opt/conda/lib/python3.7/site-packages (from absl-py->jax) (1.16.0)

WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: <https://pip.pypa.io/warnings/venv>

In [2]:

```
def obj_size_fmt(num):
    if num<10**3:
        return "{:.2f}{}".format(num, "B")
    elif ((num>=10**3)&(num<10**6)):
        return "{:.2f}K".format(num/(1024*10**3))
```

```

        return "{:.2f}{}".format(num/(1.024*10**9), "GB")
    elif ((num>=10**6)&(num<10**9)):
        return "{:.2f}{}".format(num/(1.024*10**6), "MB")
    else:
        return "{:.2f}{}".format(num/(1.024*10**9), "GB")

```

```

def memory_usage():
    memory_usage_by_variable=pd.DataFrame({k:sys.getsizeof(v)\
    for (k,v) in globals().items()} ,index=['Size'])
    memory_usage_by_variable=memory_usage_by_variable.T
    memory_usage_by_variable=memory_usage_by_variable.sort_values(by='Size',ascending=False).head(10)
    memory_usage_by_variable['Size']=memory_usage_by_variable['Size'].apply(lambda x: obj_size_fmt(x))
    return memory_usage_by_variable

```

In [3]:

#Read in customer data

```
customers = pd.read_csv('/kaggle/input/h-and-m-personalized-fashion-recommendations/customers.csv')
```

#Fill in columns with missing values and convert strings to weights

```
customers.loc[customers['FN'].isnull(), 'FN'] = float(0.0)
```

```
customers.loc[customers['Active'].isnull(), 'Active'] = float(0.0)
```

```
customers.loc[customers['fashion_news_frequency']=='NONE', 'fashion_news_frequency'] = float(0.0)
```

```
customers.loc[customers['fashion_news_frequency']=='None', 'fashion_news_frequency'] = float(0.0)
```

```
customers.loc[customers['fashion_news_frequency'].isnull(), 'fashion_news_frequency'] = float(0.0)
```

```
customers.loc[customers['fashion_news_frequency']=='Regularly', 'fashion_news_frequency'] = float(0.5)
```

```
customers.loc[customers['fashion_news_frequency']=='Monthly', 'fashion_news_frequency'] = float(1.0)
```

```
customers['fashion_news_frequency'] = customers['fashion_news_frequency'].astype('float64')
```

```

customers.loc[customers['club_member_status']=='LEFT CLUB', 'club_member_status'] = float(0.0)
customers.loc[customers['club_member_status'].isnull(), 'club_member_status'] = float(0.0)
customers.loc[customers['club_member_status']=='PRE-CREATE', 'club_member_status'] = float(0.5)
customers.loc[customers['club_member_status']=='ACTIVE', 'club_member_status'] = float(1.0)
customers['club_member_status'] = customers['club_member_status'].astype('float64')

#Assume average age if none given
average_age = customers['age'].mean()
customers.loc[customers['age'].isnull(), 'age'] = average_age

#Remove postal code
customers.drop(columns='postal_code', axis = 1, inplace = True)

customers['customer_id'] = customers['customer_id'].astype('string')

#Display results
customers.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1371980 entries, 0 to 1371979
Data columns (total 6 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   customer_id          1371980 non-null string
 1   FN                   1371980 non-null float64
 2   Active               1371980 non-null float64
 3   club_member_status   1371980 non-null float64
 4   fashion_news_frequency 1371980 non-null float64
 5   age                  1371980 non-null float64
dtypes: float64(5), string(1)
memory usage: 62.8 MB

```

In [4]:

#read in transactions and categorize years and months

```
transactions = pd.read_csv('/kaggle/input/h-and-m-personalized-fashion-recommendations/transactions_train.csv')
transactions['year'] = pd.DatetimeIndex(transactions['t_dat']).year
transactions['month'] = pd.DatetimeIndex(transactions['t_dat']).month
transactions = transactions[transactions['year'] >= 2020]
transactions['t_dat'] = transactions['t_dat'].astype('string')
transactions['customer_id'] = transactions['customer_id'].astype('string')
transactions.info()
transactions.head()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10980132 entries, 20808192 to 31788323
Data columns (total 7 columns):
#   Column          Dtype
---  -
0   t_dat           string
1   customer_id     string
2   article_id      int64
3   price           float64
4   sales_channel_id int64
5   year            int64
6   month           int64
dtypes: float64(1), int64(4), string(2)
memory usage: 670.2 MB
```

Out[4]:

	t_dat	customer_id	article_id	price	sales_channel_id	year	month
20808192	2020-01-01	0034b3dced3e565a13138bdfb5117e7321fea65388b308	835217001	0.033881	2	2020	1

20808192	2020-01-01	00410b91d62eefa76958fa5cac12f5daa7cfc0556e417d...	802930002	0.067780	2	2020	1
20808193	2020-01-01	00410b91d62eefa76958fa5cac12f5daa7cfc0556e417d...	802930002	0.067780	2	2020	1
20808194	2020-01-01	00410b91d62eefa76958fa5cac12f5daa7cfc0556e417d...	760084008	0.025407	2	2020	1
20808195	2020-01-01	004b0fb384bcab2f8e1059dd5ca68c17580365ab95c05a...	804662002	0.033881	2	2020	1
20808196	2020-01-01	004b0fb384bcab2f8e1059dd5ca68c17580365ab95c05a...	801554002	0.016932	2	2020	1

In [5]:

```
#read in articles
from sklearn.preprocessing import LabelEncoder
articles = pd.read_csv('/kaggle/input/h-and-m-personalized-fashion-recommendations/articles.csv')
labelencoder = LabelEncoder()
#convert letters to numbers for groups
articles['product_group_name'] = labelencoder.fit_transform(articles['product_group_name'])
articles['index_code'] = labelencoder.fit_transform(articles['index_code'])
#select only numeric groups
articles = articles[['article_id', 'product_code', 'product_type_no', 'product_group_name', 'graphical_appearance_no', 'colour_group_code', 'perceived_colour_value_id', 'perceived_colour_master_id', 'department_no', 'index_code', 'index_group_no', 'section_no', 'garment_group_no']]
articles.info()
articles.head(10)
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 105542 entries, 0 to 105541
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	article_id	105542 non-null	int64
1	product_code	105542 non-null	int64
2	product_type_no	105542 non-null	int64

```
3  product_group_name      105542 non-null  int64
4  graphical_appearance_no  105542 non-null  int64
5  colour_group_code       105542 non-null  int64
6  perceived_colour_value_id 105542 non-null  int64
7  perceived_colour_master_id 105542 non-null  int64
8  department_no          105542 non-null  int64
9  index_code              105542 non-null  int64
10 index_group_no          105542 non-null  int64
11 section_no              105542 non-null  int64
12 garment_group_no        105542 non-null  int64
dtypes: int64(13)
memory usage: 10.5 MB
```

Out[5]:

	article_id	product_code	product_type_no	product_group_name	graphical_appearance_no	colour_group_code	perceived_colour_value_id	perceived_colour_master.
0	108775015	108775	253	7	1010016	9	4	5
1	108775044	108775	253	7	1010016	10	3	9
2	108775051	108775	253	7	1010017	11	1	9
3	110065001	110065	306	16	1010016	9	4	5
4	110065002	110065	306	16	1010016	10	3	9
5	110065011	110065	306	16	1010016	12	1	11
6	111565001	111565	304	13	1010016	9	4	5
7	111565003	111565	302	13	1010016	13	2	11
8	111586001	111586	273	6	1010016	9	4	5
9	111593001	111593	304	13	1010016	9	4	5



In [6]:

```
#data = pd.merge(customers, transactions, how="left", on=["customer_id", "customer_id"])
```

```
#data = pd.merge(data, articles, how="left", on=["article_id", "article_id"])
#data = customers.merge(transactions, how='left', left_on='customer_id', right_on='customer_id').merge(articles, how='left', left_on='article_id', right_on='article_id')
data = transactions.merge(customers,on='customer_id').merge(articles,on='article_id')
data['ID'] = data.index
del transactions
#del customers
#saving customers for later to merge onto predictions
del articles
gc.collect()
memory_usage()
```

Out[6]:

	Size
data	4.07GB
customers	215.71MB
_iii	1.69KB
_i3	1.69KB
_4	1.18KB
—	1.18KB
_5	1.17KB
—	1.17KB
LabelEncoder	1.04KB
_i2	820.00B

In [7]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 10980132 entries, 0 to 10980131
```

```
Data columns (total 25 columns):
```

#	Column	Dtype
0	t_dat	string
1	customer_id	string
2	article_id	int64
3	price	float64
4	sales_channel_id	int64
5	year	int64
6	month	int64
7	FN	float64
8	Active	float64
9	club_member_status	float64
10	fashion_news_frequency	float64
11	age	float64
12	product_code	int64
13	product_type_no	int64
14	product_group_name	int64
15	graphical_appearance_no	int64
16	colour_group_code	int64
17	perceived_colour_value_id	int64
18	perceived_colour_master_id	int64
19	department_no	int64
20	index_code	int64
21	index_group_no	int64
22	section_no	int64
23	garment_group_no	int64
24	ID	int64

```
dtypes: float64(6), int64(17), string(2)
```


In [8]:

```
data.head(10)
```

Out[8]:

	t_dat	customer_id	article_id	price	sales_channel_id	year	month	FN	Active	club_member_status	...
0	2020-01-01	0034b3dced3e565a43438bdfb5447e7321fea65388b398...	835247001	0.033881	2	2020	1	1.0	1.0	1.0	...
1	2020-03-10	078be7d15562e689421fcad630bee1d41aea7eb518d2b1...	835247001	0.022017	1	2020	3	0.0	0.0	1.0	...
2	2020-02-04	0c30c91b3272fc36b172d7b56636fd8ce54af75a3e9368...	835247001	0.022492	2	2020	2	1.0	1.0	1.0	...
3	2020-01-01	170126f1d3345fa450b87b147906643a3659ccb6ef2871...	835247001	0.031508	2	2020	1	0.0	0.0	1.0	...
4	2020-01-01	170126f1d3345fa450b87b147906643a3659ccb6ef2871...	835247001	0.031576	2	2020	1	0.0	0.0	1.0	...
5	2020-02-11	2414b5f86a1742ebbadc781561ddfac6ca0177a9ef08f1...	835247001	0.027441	2	2020	2	0.0	0.0	1.0	...
6	2020-03-15	31b07944e12db276190a729edce320718ecc3a67121e78...	835247001	0.030492	2	2020	3	1.0	1.0	1.0	...
7	2020-01-01	33992d5bc04fbf1ddac7df567b919ac37d814f293013af...	835247001	0.032068	2	2020	1	0.0	0.0	1.0	...
8	2020-01-01	4444f737c7bc41faef68d30c26957d1c41bdeaef36e926...	835247001	0.033881	2	2020	1	0.0	0.0	1.0	...
9	2020-01-01	459913ba6177f175e6c76a1367f839008ccbe5e46756f5...	835247001	0.033881	2	2020	1	0.0	0.0	1.0	...

10 rows × 25 columns

In [9]:

```
#eliminate null values while maintaining customer count
#data.loc[data['t_dat'].isnull(),'t_dat'] = '00-00-0000'
#data.loc[data['article_id'].isnull(),'article_id'] = float(0.0)
#data.loc[data['price'].isnull(),'price'] = float(0.0)
#data.loc[data['sales_channel_id'].isnull(),'sales_channel_id'] = float(0.0)
#data.loc[data['year'].isnull(),'year'] = float(0.0)
#data.loc[data['month'].isnull(),'month'] = float(0.0)
#data.loc[data['product_code'].isnull(),'product_code'] = float(0.0)
#data.loc[data['product_type_no'].isnull(),'product_type_no'] = float(0.0)
#data.loc[data['product_group_name'].isnull(),'product_group_name'] = float(0.0)
#data.loc[data['graphical_appearance_no'].isnull(),'graphical_appearance_no'] = float(0.0)
#data.loc[data['colour_group_code'].isnull(),'colour_group_code'] = float(0.0)
#data.loc[data['perceived_colour_value_id'].isnull(),'perceived_colour_value_id'] = float(0.0)
#data.loc[data['perceived_colour_master_id'].isnull(),'perceived_colour_master_id'] = float(0.0)
#data.loc[data['department_no'].isnull(),'department_no'] = float(0.0)
#data.loc[data['index_code'].isnull(),'index_code'] = float(0.0)
#data.loc[data['index_group_no'].isnull(),'index_group_no'] = float(0.0)
#data.loc[data['section_no'].isnull(),'section_no'] = float(0.0)
#data.loc[data['garment_group_no'].isnull(),'garment_group_no'] = float(0.0)
#Turns out, this was a bad idea. The model predicts 0 all the time and becomes horribly inaccurate.

data.isnull().sum()
```

Out[9]:

t_dat	0
customer_id	0
article_id	0
price	0
sales_channel_id	0
year	0

month	0
FN	0
Active	0
club_member_status	0
fashion_news_frequency	0
age	0
product_code	0
product_type_no	0
product_group_name	0
graphical_appearance_no	0
colour_group_code	0
perceived_colour_value_id	0
perceived_colour_master_id	0
department_no	0
index_code	0
index_group_no	0
section_no	0
garment_group_no	0
ID	0

dtype: int64

In [10]:

```
#JAX FOR RANDOM SAMPLING
import jax.numpy as jnp
from jax import random
#tried using jax_sampling methods and ran into many issues, switched to jax random generation

#Using JAX to create a random array of indexes and use that to take a sample out of the data to pass to X
seed = 1234
key = jax.random.PRNGKey(seed)
#SAMPLE SIZE
#continued for random sampling
```

```

#optimized for accuracy vs time
shape = (25000,)
index_rand = jax.random.randint(key, shape, 0, len(data)-1)
index_rand = np.array(index_rand)

#JAX version
x = data.iloc[index_rand, :].copy()
y = x['article_id']
#Dropping non-numerical data for later
#going to rejoin data to data_num based on ID
x.drop(columns=['article_id', 't_dat', 'customer_id'], axis = 1, inplace = True)
indexes = data[['article_id', 't_dat', 'customer_id', 'ID']]
data.drop(columns=['article_id', 't_dat', 'customer_id'], axis = 1, inplace = True)

```

In [11]:

```

# split the dataset into the training set and test set
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state = 0)
del x
del y
gc.collect()
memory_usage()

```

Out[11]:

	Size
indexes	2.27GB
data	1.97GB
customers	215.71MB
x_train	3.59MB
y_test	808.47KB

x_test	898.47KB
y_train	312.53KB
index_rand	97.76KB
y_test	78.16KB
_8	3.74KB
—	3.74KB

In [12]:

```
from sklearn.neighbors import KNeighborsClassifier
x_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20000 entries, 447385 to 4698435
Data columns (total 22 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   price                                20000 non-null  float64
 1   sales_channel_id                     20000 non-null  int64
 2   year                                 20000 non-null  int64
 3   month                                20000 non-null  int64
 4   FN                                    20000 non-null  float64
 5   Active                               20000 non-null  float64
 6   club_member_status                   20000 non-null  float64
 7   fashion_news_frequency               20000 non-null  float64
 8   age                                   20000 non-null  float64
 9   product_code                         20000 non-null  int64
10   product_type_no                      20000 non-null  int64
11   product_group_name                   20000 non-null  int64
12   graphical_appearance_no              20000 non-null  int64
```

13	colour_group_code	20000	non-null	int64
14	perceived_colour_value_id	20000	non-null	int64
15	perceived_colour_master_id	20000	non-null	int64
16	department_no	20000	non-null	int64
17	index_code	20000	non-null	int64
18	index_group_no	20000	non-null	int64
19	section_no	20000	non-null	int64
20	garment_group_no	20000	non-null	int64
21	ID	20000	non-null	int64

dtypes: float64(6), int64(16)

memory usage: 3.5 MB

In [13]:

```
k = 1
memory_usage()
#Train Model and Predict
neighbors = KNeighborsClassifier(n_neighbors = k).fit(x_train,y_train)
```

In [14]:

```
yhat = neighbors.predict(x_test)
yhat[0:5]
```

Out[14]:

```
array([810227004, 636323002, 866383001, 806731001, 824337001])
```

In [15]:

```
from sklearn import metrics
```

```
print("Train set Accuracy: ", metrics.accuracy_score(y_train, neighbors.predict(x_train)))
print("Test set Accuracy: ", metrics.accuracy_score(y_test, yhat))
```

```
Train set Accuracy:  1.0
Test set Accuracy:  0.6912
```

In [16]:

```
x_train.info()
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 20000 entries, 447385 to 4698435
```

```
Data columns (total 22 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	price	20000 non-null	float64
1	sales_channel_id	20000 non-null	int64
2	year	20000 non-null	int64
3	month	20000 non-null	int64
4	FN	20000 non-null	float64
5	Active	20000 non-null	float64
6	club_member_status	20000 non-null	float64
7	fashion_news_frequency	20000 non-null	float64
8	age	20000 non-null	float64
9	product_code	20000 non-null	int64
10	product_type_no	20000 non-null	int64
11	product_group_name	20000 non-null	int64
12	graphical_appearance_no	20000 non-null	int64
13	colour_group_code	20000 non-null	int64

14	perceived_colour_value_id	20000	non-null	int64
15	perceived_colour_master_id	20000	non-null	int64
16	department_no	20000	non-null	int64
17	index_code	20000	non-null	int64
18	index_group_no	20000	non-null	int64
19	section_no	20000	non-null	int64
20	garment_group_no	20000	non-null	int64
21	ID	20000	non-null	int64

dtypes: float64(6), int64(16)

memory usage: 3.5 MB

<class 'pandas.core.frame.DataFrame'>

Int64Index: 10980132 entries, 0 to 10980131

Data columns (total 22 columns):

#	Column	Dtype
---	-----	-----
0	price	float64
1	sales_channel_id	int64
2	year	int64
3	month	int64
4	FN	float64
5	Active	float64
6	club_member_status	float64
7	fashion_news_frequency	float64
8	age	float64
9	product_code	int64
10	product_type_no	int64
11	product_group_name	int64
12	graphical_appearance_no	int64
13	colour_group_code	int64
14	perceived_colour_value_id	int64
15	perceived_colour_master_id	int64
16	department_no	int64
17	index_code	int64


```
18  index_group_no          int64
19  section_no              int64
20  garment_group_no        int64
21  ID                      int64
```

```
dtypes: float64(6), int64(16)
```

```
memory usage: 1.9 GB
```

In [17]:

```
result_data = neighbors.predict(data)
```

In [18]:

```
len(result_data)
len(data)
```

Out[18]:

```
10980132
```

In [19]:

```
data = data.merge(indexes, on='ID')
data['article'] = result_data.tolist()
```

In [20]:

```
data_ndg = data.drop_duplicates(subset = ['customer_id', 'article'])
data_ngd = data_ndg.groupby('customer_id')
data_ndg.head(10)
```

```
data_ndg.value_counts()
```

Out[20]:

	price	sales_channel_id	year	month	FN	Active	club_member_status	fashion_news_frequency	age	product_code	...	department_no	index_code	ir
0	0.033881	2	2020	1	1.0	1.0	1.0	0.5	42.0	835247	...	1322	0	1
1	0.022017	1	2020	3	0.0	0.0	1.0	0.0	22.0	835247	...	1322	0	1
2	0.022492	2	2020	2	1.0	1.0	1.0	0.5	45.0	835247	...	1322	0	1
3	0.031508	2	2020	1	0.0	0.0	1.0	0.0	37.0	835247	...	1322	0	1
5	0.027441	2	2020	2	0.0	0.0	1.0	0.0	25.0	835247	...	1322	0	1
6	0.030492	2	2020	3	1.0	1.0	1.0	0.5	42.0	835247	...	1322	0	1
7	0.032068	2	2020	1	0.0	0.0	1.0	0.0	29.0	835247	...	1322	0	1
8	0.033881	2	2020	1	0.0	0.0	1.0	0.0	49.0	835247	...	1322	0	1
9	0.033881	2	2020	1	0.0	0.0	1.0	0.0	37.0	835247	...	1322	0	1
11	0.030492	2	2020	1	1.0	1.0	1.0	0.5	38.0	835247	...	1322	0	1

10 rows × 26 columns



In [21]:

```
data_ndg['customer_id'].value_counts()
```

Out[21]:

b637a3e7d8b0caa947aaefd609b8d84a9ee962cf0a52a51bac507ffc2bf1b741 539
be1981ab818cf4ef6765b2ecaea7a2cbf14ccd6e8a7ee985513d9e8e53c6d91b 499
cd04ec2726dd58a8c753e0d6423e57716fd9ebcf2f14ed6012e7e5bea016b4d6 442
863f0e03da282ae32a76775ce55d8a4605a85c84a26066e1ad0e9469e8c40e68 412
a65f77281a528bf5c1e9f270141d601d116e1df33bf9df512f495ee06647a9cc 405
...
174d69cb658377048e5d730b1fef62cda4024b719debf7b7b03735ee027f13e7 1
a4b158c57d4483d1cb87f515fe2742bbfddff3128133d45ddb45365e6b5052a0d 1
3372a6bd1f2b06e339cd7c4ac0c3cad4d24bb5a0430c52e76a7511349d6ec6d1 1
bbe5bbdc885e15b232c6d19a1d4524c448cff17942224108ed664c0b193ac054 1
daaceb3e5ec42f538f20dc11a0baccea645c4e6dd22088988bd9289fa0715881 1
Name: customer_id, Length: 862724, dtype: Int64

In [22]:

```
data_ndg.head(10)
```

Out[22]:

	price	sales_channel_id	year	month	FN	Active	club_member_status	fashion_news_frequency	age	product_code	...	department_no	index_code	ir
0	0.033881	2	2020	1	1.0	1.0	1.0	0.5	42.0	835247	...	1322	0	1
1	0.022017	1	2020	3	0.0	0.0	1.0	0.0	22.0	835247	...	1322	0	1
2	0.022492	2	2020	2	1.0	1.0	1.0	0.5	45.0	835247	...	1322	0	1
3	0.031508	2	2020	1	0.0	0.0	1.0	0.0	37.0	835247	...	1322	0	1
5	0.027441	2	2020	2	0.0	0.0	1.0	0.0	25.0	835247	...	1322	0	1
6	0.030492	2	2020	3	1.0	1.0	1.0	0.5	42.0	835247	...	1322	0	1

7	price 0.032068	sales_channel_id 2	year 2020	month 1	FN 0.0	Active 0.0	club_member_status 1.0	fashion_news_frequency 0.0	age 29.0	product_code 835247	...	department_no 1322	index_code 0	ir
8	0.033881	2	2020	1	0.0	0.0	1.0	0.0	49.0	835247	...	1322	0	1
9	0.033881	2	2020	1	0.0	0.0	1.0	0.0	37.0	835247	...	1322	0	1
11	0.030492	2	2020	1	1.0	1.0	1.0	0.5	38.0	835247	...	1322	0	1

10 rows × 26 columns



In [23]:

```
result_data=data_ndg[['customer_id','article', 'ID']]
result_data
```

Out[23]:

	customer_id	article	ID
0	0034b3dced3e565a43438bdfb5447e7321fea65388b398...	821673001	0
1	078be7d15562e689421fcad630bee1d41aea7eb518d2b1...	821673001	1
2	0c30c91b3272fc36b172d7b56636fd8ce54af75a3e9368...	821673001	2
3	170126f1d3345fa450b87b147906643a3659ccb6ef2871...	821673001	3
5	2414b5f86a1742ebbadc781561ddf6ca0177a9ef08f1...	821673001	5
...
10980126	91539e75dac101c6c1761a08842f8fbe34d1e15f58abe3...	646829002	10980126
10980127	cf078f032736773511094c34b76010dee62510152ddeb...	907317001	10980127
10980128	eaf28084976ef4eaf6b3be9c2a91ca6679e91c74959bb5...	907317001	10980128
10980129	ec794d6268ee3c75b5bbea014e3299d994bcd1180ecf6d...	850134003	10980129
10980130	ec794d6268ee3c75b5bbea014e3299d994bcd1180ecf6d...	547365012	10980130

In [24]:

```
result_data['customer_id'].value_counts()
```

Out[24]:

b637a3e7d8b0caa947aaefd609b8d84a9ee962cf0a52a51bac507ffc2bf1b741	539
be1981ab818cf4ef6765b2ecaea7a2cbf14ccd6e8a7ee985513d9e8e53c6d91b	499
cd04ec2726dd58a8c753e0d6423e57716fd9ebcf2f14ed6012e7e5bea016b4d6	442
863f0e03da282ae32a76775ce55d8a4605a85c84a26066e1ad0e9469e8c40e68	412
a65f77281a528bf5c1e9f270141d601d116e1df33bf9df512f495ee06647a9cc	405
...	
174d69cb658377048e5d730b1fef62cda4024b719debfb7b03735ee027f13e7	1
a4b158c57d4483d1cb87f515fe2742bbfdff3128133d45ddb45365e6b5052a0d	1
3372a6bd1f2b06e339cd7c4ac0c3cad4d24bb5a0430c52e76a7511349d6ec6d1	1
bbe5bbdc885e15b232c6d19a1d4524c448cff17942224108ed664c0b193ac054	1
daaceb3e5ec42f538f20dc11a0baccea645c4e6dd22088988bd9289fa0715881	1
Name: customer_id, Length: 862724, dtype: Int64	

In [25]:

```
result_final = result_data.groupby('customer_id').sum().reset_index()
result_final = result_final.merge(customers,how='right', on='customer_id')
result_final = result_final[['customer_id', 'article']]
```

Out[25]:

customer_id	0
-------------	---

article 509256
dtype: int64

In [26]:

```
result_final.columns = ['customer_id','prediction']  
#Guess most popular product for every other customer with no transaction history  
most_popular_product = result_final['prediction'].mode()  
result_final.loc[result_final['prediction'].isnull(),'prediction'] = int(most_popular_product)  
  
#pad with 0 and convert to string for article_ids  
result_final['prediction'] = result_final['prediction'].astype('int')  
result_final['prediction'] = result_final['prediction'].astype('string')  
result_final['prediction'] = result_final['prediction'].apply(lambda x: x.zfill(10))  
  
result_final['prediction'] = result_final['prediction'].astype('string')  
result_final['customer_id'] = result_final['customer_id'].astype('string')  
result_final.head(20)
```

Out[26]:

	customer_id	prediction
0	00000dbacae5abe5e23885899a1fa44253a17956c6d1c3...	4842791057
1	0000423b00ade91418cceaf3b26c6af3dd342b51fd051e...	16928243327
2	000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...	8352360041
3	00005ca1c9ed5f5146b52ac8639a40ca9d57aeff4d1bd2...	0706016001
4	00006413d8573cd20ed7128e53b7b13819fe5cfc2d801f...	5550285086
5	000064249685c11552da43ef22a5030f35a147f723d5b0...	0706016001
6	0000757967448a6cb83efb3ea7a3fb9d418ac7adf2379d...	1168039017
7	00007d2de826758b65a93dd24ce629ed66842531df6699...	4879841163
8	00007e8d4e54114b5b2a9b51586325a8d0fa74ea23ef77...	1670037002
9	00008169a21b50b3d117c97135a25b1201a8c58997f787	0706016001

9	0000400a21b30b3d147007133c23b4201a0c000071707...	0700010001
10	customer_id 0000945f66de1a11d9447609b8b41b1bc987ba185a5496...	prediction 1520168016
11	000097d91384a0c14893c09ed047a963c4fc6a5c021044...	0706016001
12	00009c2aeae8761f738e4f937d9be6b49861a66339c2b1...	1711185002
13	00009d946eec3ea54add5ba56d5210ea898def4b46c685...	38705401704
14	0000ae1bbb25e04bdc7e35f718e852adfb3fbb72ef38b3...	0706016001
15	0000b2f1829e23b24feec422ef13df3ccedaedc85368e6...	11697524211
16	0000b7a134c3ec0d8842fad1fd4ca28517424c14fc4848...	0706016001
17	0000b95f630aaa9313028ce9c41154bb95ac7afa34f55b...	0706016001
18	0000c97821eb48d0e590fd309133f0a6c08f7750f64ccc...	4868755058
19	0000d6c053fc8f9389d4565051f12402d5774aa4a9d2e5...	0706016001

In [28]:

```
#check for null values
result_final.isnull().sum()
```

Out[28]:

```
customer_id      0
prediction        0
dtype: int64
```

In [27]:

```
import reprlib
submit = result_final.to_csv(index=False)
#limit output displayed
print(reprlib.repr(submit))
```

'customer_id,...,0706016001\n'